

Janvijay Singh

Champaign, Illinois

 <https://iamjanvijay.github.io/>

 jvsingh@illinois.edu  +1 (470) 753-5241

EDUCATION

- **University of Illinois Urbana-Champaign** Urbana, IL
Ph.D. in Computer Science · Focus: Reasoning in LLMs · Advisors: Profs. Dilek Hakkani-Tur & Gokhan Tur Aug 2024 - Present
- **Georgia Institute of Technology** Atlanta, GA
M.S. in Computer Science · Specialization: Machine Learning · GPA: 3.91/4.00 Aug 2021 - May 2023
- **Indian Institute of Technology Varanasi** Varanasi, India
B.Tech. in Computer Science and Engineering · GPA: 9.62/10.00 · Department Rank: 2/65 July 2014 - May 2018

INTERESTS & SKILLS

- **Research Interests:** Reasoning in Large Language Models (LLMs), with a focus on characterizing and modeling design-attributes of human-like reasoning that are missing in LLMs. *Keywords:* Abstractions, Systematic Generalization (Easy-to-Hard, Weak-to-Strong), Continual Learning, Faithfulness, and Interpretability for more human-like reasoning.
- **Languages:** Python, Bash, C++, C, Java, Objective-C.
- **Technologies:** PyTorch, TensorFlow, MXNet, JAX, CUDA, HuggingFace, DeepSpeed, Megatron-LM, vLLM, TGI, NVIDIA Triton, TensorBoard, Numpy, Pandas, Scikit-Learn, Matplotlib, MATLAB, Docker, Kubernetes, Helm, Terraform, GCP, AWS, FastAPI, gRPC, Django, Flask, Git, LATEX.

PUBLICATIONS

- JANVIJAY SINGH, AUSTIN XU, YILUN ZHOU, YEFAN ZHOU, DILEK HAKKANI-TUR, SHAFIQ JOTY. "On the Shelf Life of Finetuned LLM-Judges: Future Proofing, Backward Compatibility, and Question Generalization.". Under review at International Conference on Learning Representations (ICLR) 2026, Preprint. <https://arxiv.org/pdf/2509.23542.pdf>
- YEFAN ZHOU, AUSTIN XU, YILUN ZHOU, JANVIJAY SINGH, JIANG GUI, SHAFIQ JOTY. "Variation in verification: Understanding verification dynamics in large language models". Under review at International Conference on Learning Representations (ICLR) 2026, Preprint. <https://arxiv.org/pdf/2509.17995.pdf>
- JANVIJAY SINGH, VILEM ZOUHAR, MRINMAYA SACHAN. "Enhancing Textbooks with Visuals from the Web for Improved Learning". In Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP) 2023. <https://aclanthology.org/2023.emnlp-main.731.pdf>
- JANVIJAY SINGH⁼, MUKUND RUNGTA⁼, DIYI YANG, SAIF M. MOHAMMAD. "Forgotten Knowledge: Examining the Citational Amnesia in NLP". In Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL) 2023. *Best Paper Honourable Mention*. <https://aclanthology.org/2023.acl-long.341.pdf>
- JANVIJAY SINGH, FAN BAI, ZHEN WANG. "Entity Tracking via Effective Use of Multi-Task Learning Model and Mention-guided Decoding". In Proceedings of the European Chapter of the Association for Computational Linguistics (EACL) 2023. <https://aclanthology.org/2023.eacl-main.90.pdf>
- MUKUND RUNGTA⁼, JANVIJAY SINGH⁼, SAIF M. MOHAMMAD, DIYI YANG. "Geographic Citation Gaps in NLP Research". In Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP) 2022. <https://aclanthology.org/2022.emnlp-main.89.pdf>

⁼ denotes equal contribution.

EXPERIENCE

- **Salesforce Research** Palo Alto, CA
LLM Reasoning – AI Research Intern · Manager: Drs. Austin Xu & Shafiq Joty May 2025 - Aug 2025
 - Designed and implemented systematic experiments across multiple judge-training recipes to study their generalization from weak to strong exam-taker responses.
- **TikTok** San Jose, CA
Social Recommendation – Machine Learning Engineer 2 Jan 2024 - July 2024
 - Performed feature engineering and value tree tuning for a set of neural recommendation models, significantly improving various engagement and followership metrics through A/B testing on live traffic.
- **Verneek AI** Manhattan, NY
Personalised Search for E-Commerce – Machine Learning Researcher/Engineer June 2023 - Jan 2024

- Designed a personalized multimodal retrieval system using SOTA embedding models and developed its evaluation framework, enabling text and interaction-based personalization. Fine-tuned embedding models using adversarially generated e-commerce data from LLMs (LLAMA 2 70B, GPT-4), improving personalization results.

• Walmart Group Bangalore, India
Aug 2018 - July 2021

Voice-Assistant for E-Commerce – Applied Scientist 2

- **Speech Recognition for Indic Languages:** Improved speech recognition for Indic languages using hierarchical CTC-based neural models. Achieved a ~38% relative WER improvement via joint-fusion of neural language models and Transducer loss. Improved decoder efficiency by introducing algorithmic approximations.
- **Speech Synthesis for Indic Languages:** Benchmarked SOTA architectures, including Tacotron2, WaveGlow, WaveNet, and ClariNet. Innovated a multilingual training framework using automatic transliteration and learnable sentence-style embeddings to enhance conversational prosody. Outperformed Google and Amazon APIs in code-mixed and vernacular domains.
- Used multi-node training with Torchrun & DeepSpeed; developed 2x fast production-ready CUDA/cuDNN code.

• Microsoft Hyderabad, India
May 2017 - Aug 2017

MS Excel – Software Engineer Intern

- Developed UI dialogs and back-end routines for pivot-table functionality in Microsoft Excel for MacOS. Offered a full-time role for outstanding contributions.

PROJECTS

- **RNN-Transducer Prefix Beam Search:** Optimized prefix-beam search for speech recognition model with caching, batching, and 2-D beam pruning. >10x speed-up. Studied the increment in error-rates caused by these approximations. Open-sourced CUDA implementation. https://github.com/iamjanvijay/rnnt_decoder_cuda (66 stars)
- **RNN-Transducer Loss Function:** Devised diagonal parallelism to reduce time complexity from $\mathcal{O}(T * U)$ to $\mathcal{O}(T + U)$. Open-sourced TensorFlow implementation as a Python package. <https://github.com/iamjanvijay/rnnt> (45 stars)

ACHIEVEMENTS

- Honorable Mention for Best Paper Award, ACL conference, 2023.
- ETH Zurich Summer Research Fellowship (<0.8% acceptance), 2022.
- Best Team Award at Flipkart for Text-to-Speech excellence in vernacular domain, 2021.
- Winner, FinSBD-2 Task at FinNLP@IJCAI (Prize: USD 1000), 2020.
- Runner-Up, Walmart Data Science Hackathon (Prize: INR 20,000), 2019.
- Candidate Master on Codeforces (Max Rating: 1920), 2017.
- Ranked Top 0.82% globally in Algorithms on Hackerrank, 2017.
- 67th in ACM ICPC India Regionals (out of 402 teams), 2016.

RELEVANT COURSEWORK

- **Graduate Courses:** Deep Learning*, Machine Learning*, Natural Language Processing*, Computational Data Analysis*, Machine Learning Theory, Advanced Algorithms and Uncertainty, Computational Social Science, Languages and Computers, Language Interfaces and Communication.
- **Undergraduate Courses:** Artificial Intelligence, Computer Vision, Intelligent Computing, Theory of Computation, Optimisation Techniques, Probability and Statistics, Mathematical Methods, Operation Research, Discrete Mathematics, Computer Programming and Linux*.

* indicates teaching assistantship.